Maynooth University
National University of Ireland Maynooth

DHCP, DNS & IP Address Management

Bart Busschots

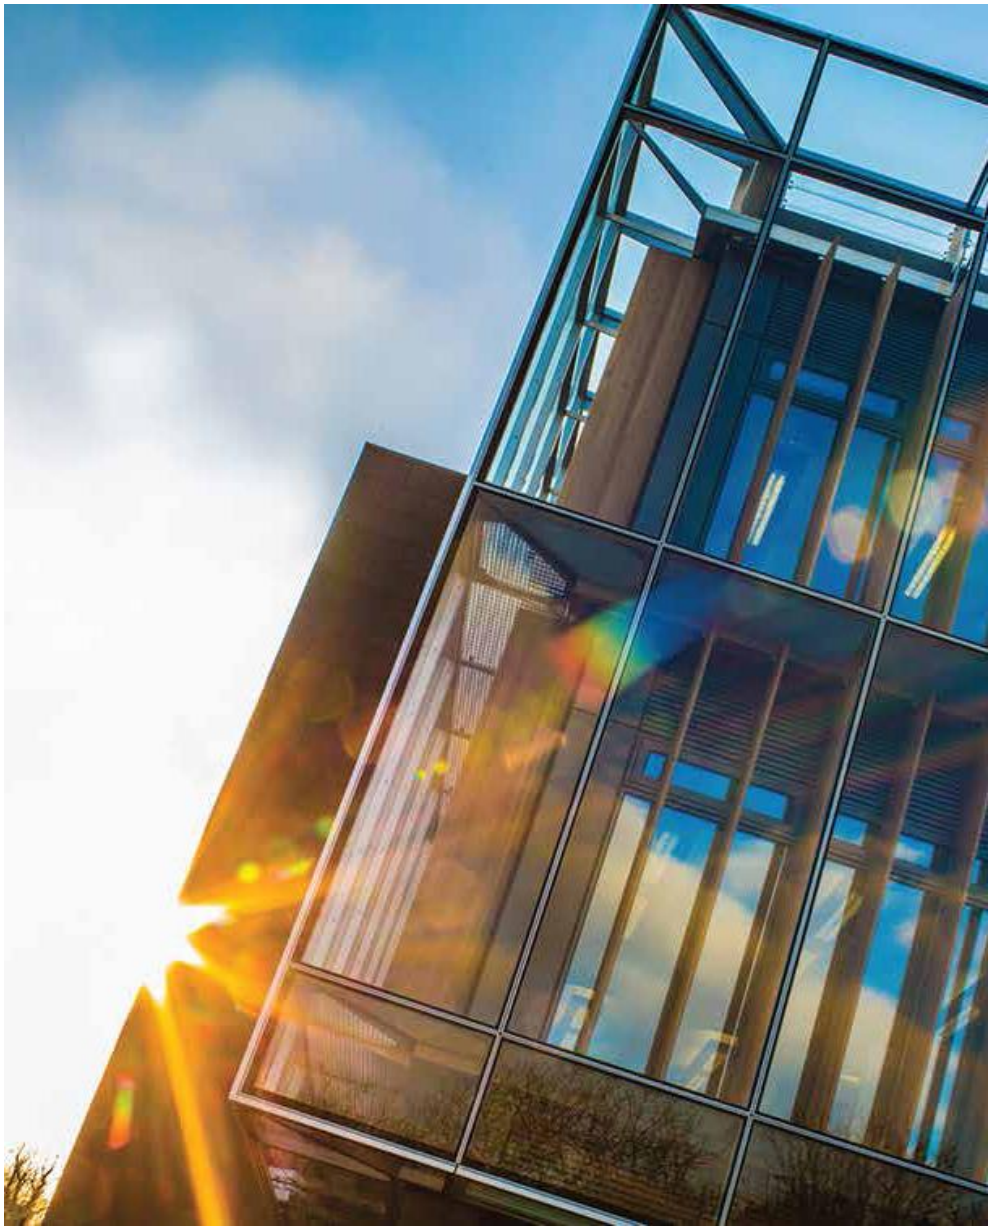

Outline

- The Back-story
  - How did we get to where we were?
  - What was the problem?
  - Why the need for change?
- Scope & Requirements
  - What must we have?
  - What would be nice to have?
- Our Solution
  - The over-all design
  - Each component in more detail
- How did it Go?
  - How did we migrate?
  - How have things changed?

Maynooth University
National University of Ireland Maynooth

The Back-story

A Long Time Ago …

- DHCP was initially deployed in the early 2000s
- Each device was assigned an IP (DHCP Reservation/Static Lease)
- A simple two-table database (‘*the host list*’) was maintained
  - A table of subnet data
  - A table of host data
- DHCP config, and Forward & Reverse DNS zones generated from *the host list*
- Single DHCP server
- Single pair of BIND DNS servers providing:
  - DNS Resolution for clients on campus
  - Hosting of public and private forward zones (using views)
  - Hosting of pubic and private reverse zones (using views)

The Cracks Begin to Show …

Implicit Assumptions

- Computers don't move around
- One subnet per department, and computers don't move between departments
- One-to-one-to-one mapping between IP addresses, MAC addresses, and hostnames
- Computers need static IPs & DNS names for peer-to-peer sharing (shared USB printers & local file shares)
- Only devices purchased by the University connect to the network

Modern Reality

- Staff carry devices with them from building to building all the time
- Many staff work for/with multiple departments
- Devices often connect via multiple MAC addresses (ethernet & WiFi)
- MAC addresses can move between devices (shared ethernet dongles)
- Peer-to-peer sharing is not a requirement anymore (replaced by central file shares, Office 365 etc.)
- Staff use many personally owned devices

Chickens Come Home to Roost …

- Adding a new computers to the network was slow — calls had to be passed between sections within IT Services (User Support → Infrastructure → User Support)
- Reservations for all devices resulted in a false scarcity of IP addresses (total registered devices high, concurrent usage low)
- *Host list* could only store data for basic DNS records
  - All other DNS records were hand-coded into the zone files.
  - Usage of other kinds of DNS record increasing (particularly SRV & TXT)
- *Host list* couldn't deal with IPv6
- Becoming ever more brittle with age
  - Original developers no longer with the institution
  - Requirements drifted over time — scripts became ever more *hacky*
  - Only a small number of staff had the knowledge needed to hand-coded the DNS records not handled by *the host list*, or to deal with outages or problems

Creaks & Groans …

- DHCP Service provided by single physical server
  - No resilience
  - Outages each time the system needed to be patched
- DNS Resolution & Hosting of Authoritative Zones provided by single pair of master/slave BIND DNS Servers
  - Use of a single cluster for both hosting zones and DNS name resolution was once standard practice, but now considered insecure
  - Used BIND views which are not compatible with DNS SEC

Maynooth University
National University of Ireland Maynooth

Scope & Requirements

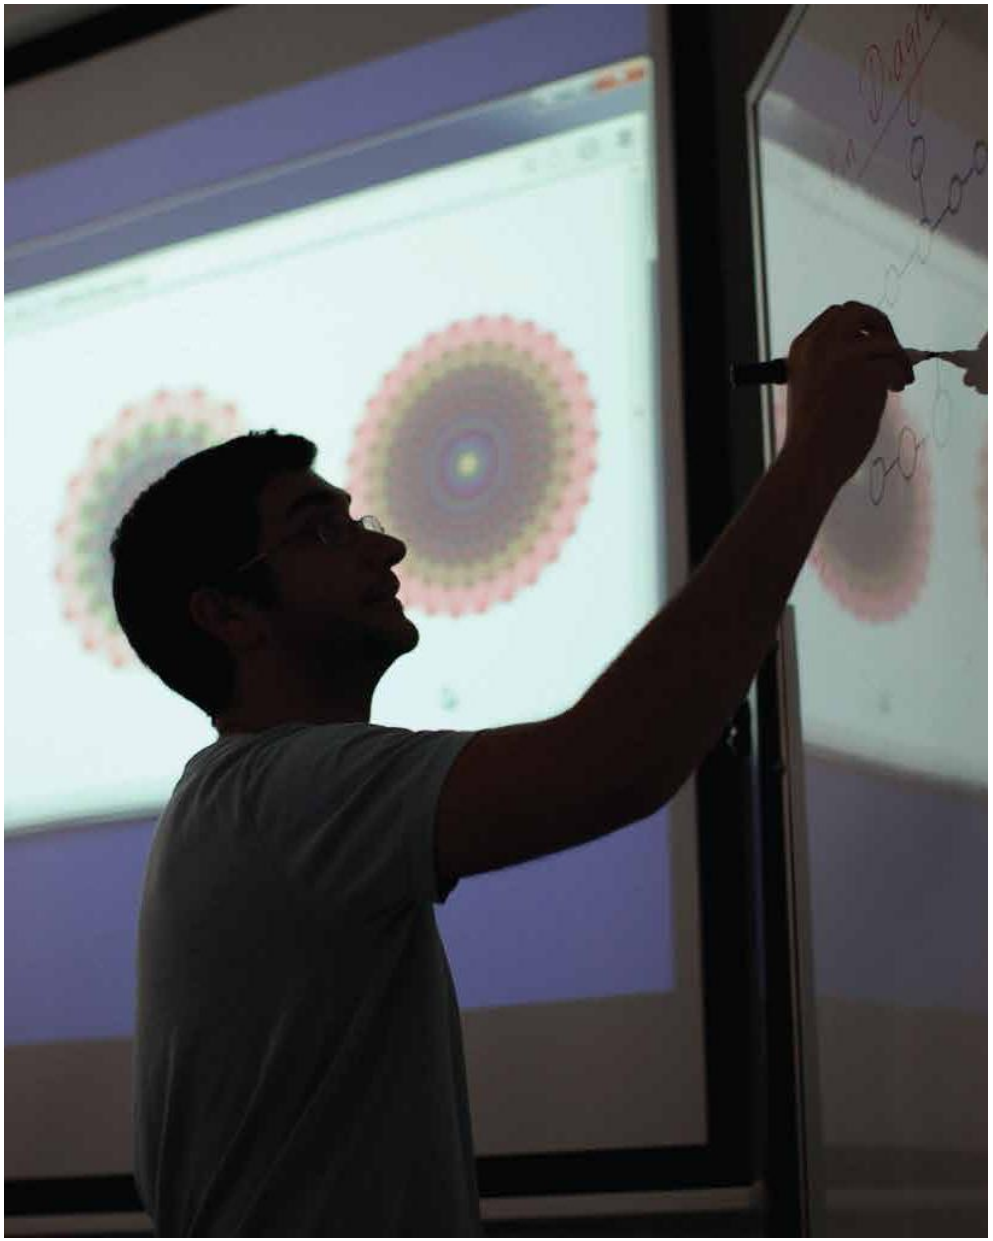

Needed

Delivered Services

- Resilient DNS name resolution
- Resilient hosting of authoritative DNS zones (public & private, and forward & reverse)
- Resilient DHCP service
- DDI management GUI for IT Services staff
- Granularly controlled access to DDI data for other stake holders, e.g. departmental technicians

Behind the Scenes

- Data stores to act as the single authoritative source for:
  - **IPv4** & **IPv6** info including **subnet definitions**, IP ranges for **DHCP pools**, **IP assignments** & per-subnet **DHCP options**
  - Authoritative **DNS records**
- The Data stores must have:
  - A well defined and documented format
  - One or more well documented APIs for querying and updating data
- Automatic generation (based on above data sources) of:
  - DHCP config
  - Reverse DNS zones

Wanted

- Low-cost— no realistic likely-hood of approval for significant budget (ruled out appliances)
- Open systems — no black boxes
  - Open source preferable
  - Well documented data storage schemes and APIs
- Widely used *'Industry Standard'* systems
- Minimise custom scripting
  - Limit custom scripts to a middle-layer between '*standard*' systems
  - Focus on implementation of MU-specific business rules
  - Avoid direct manipulation of the raw data — use well documented APIs were ever possible

Maynooth University
National University of Ireland Maynooth

Our Solution

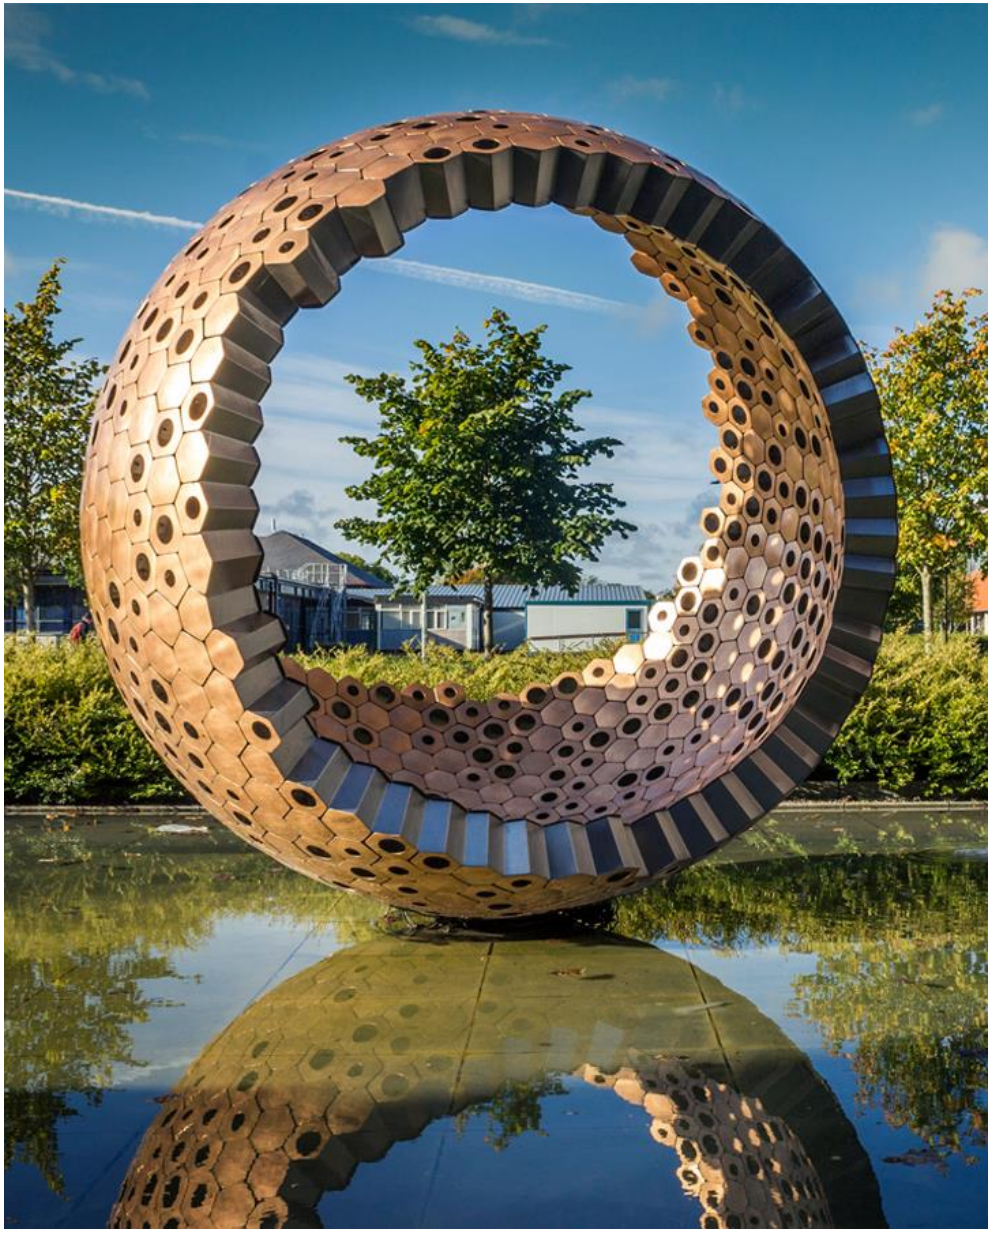

Components

Data Stores

- **php{IPAM}**
  - IPv4 & IPv6 IPAM
- **PowerDNS Server**
  - Forward DNS zones
  - *Blind Master* to Global & Local BIND clusters
- **MU DDI Config File (JSON)**
  - Config for all MU DDI Scripts

Middleware

- **MU DDI Scripts**
  - Auto-generate DHCP config
  - Auto-generate Reverse DNS Zones

Services Provided

- DHCP — active-active **ISC DHCPD** cluster
- DNS Name Resolution — **PowerDNS Recursor**
- DNS Zone Hosting — 2 **ISC BIND** clusters
  - Global (public forward & reverse domains)
  - Local (private forward domains and private version of reverse domains)

Admin GUIs

- **php{IPAM}**
- **NSEdit**
- **MySQL Client of Choice**

MU DDI Overview (DHCP, DNS & IPAM)

Bart Busschots - November 2017 (v1.0)

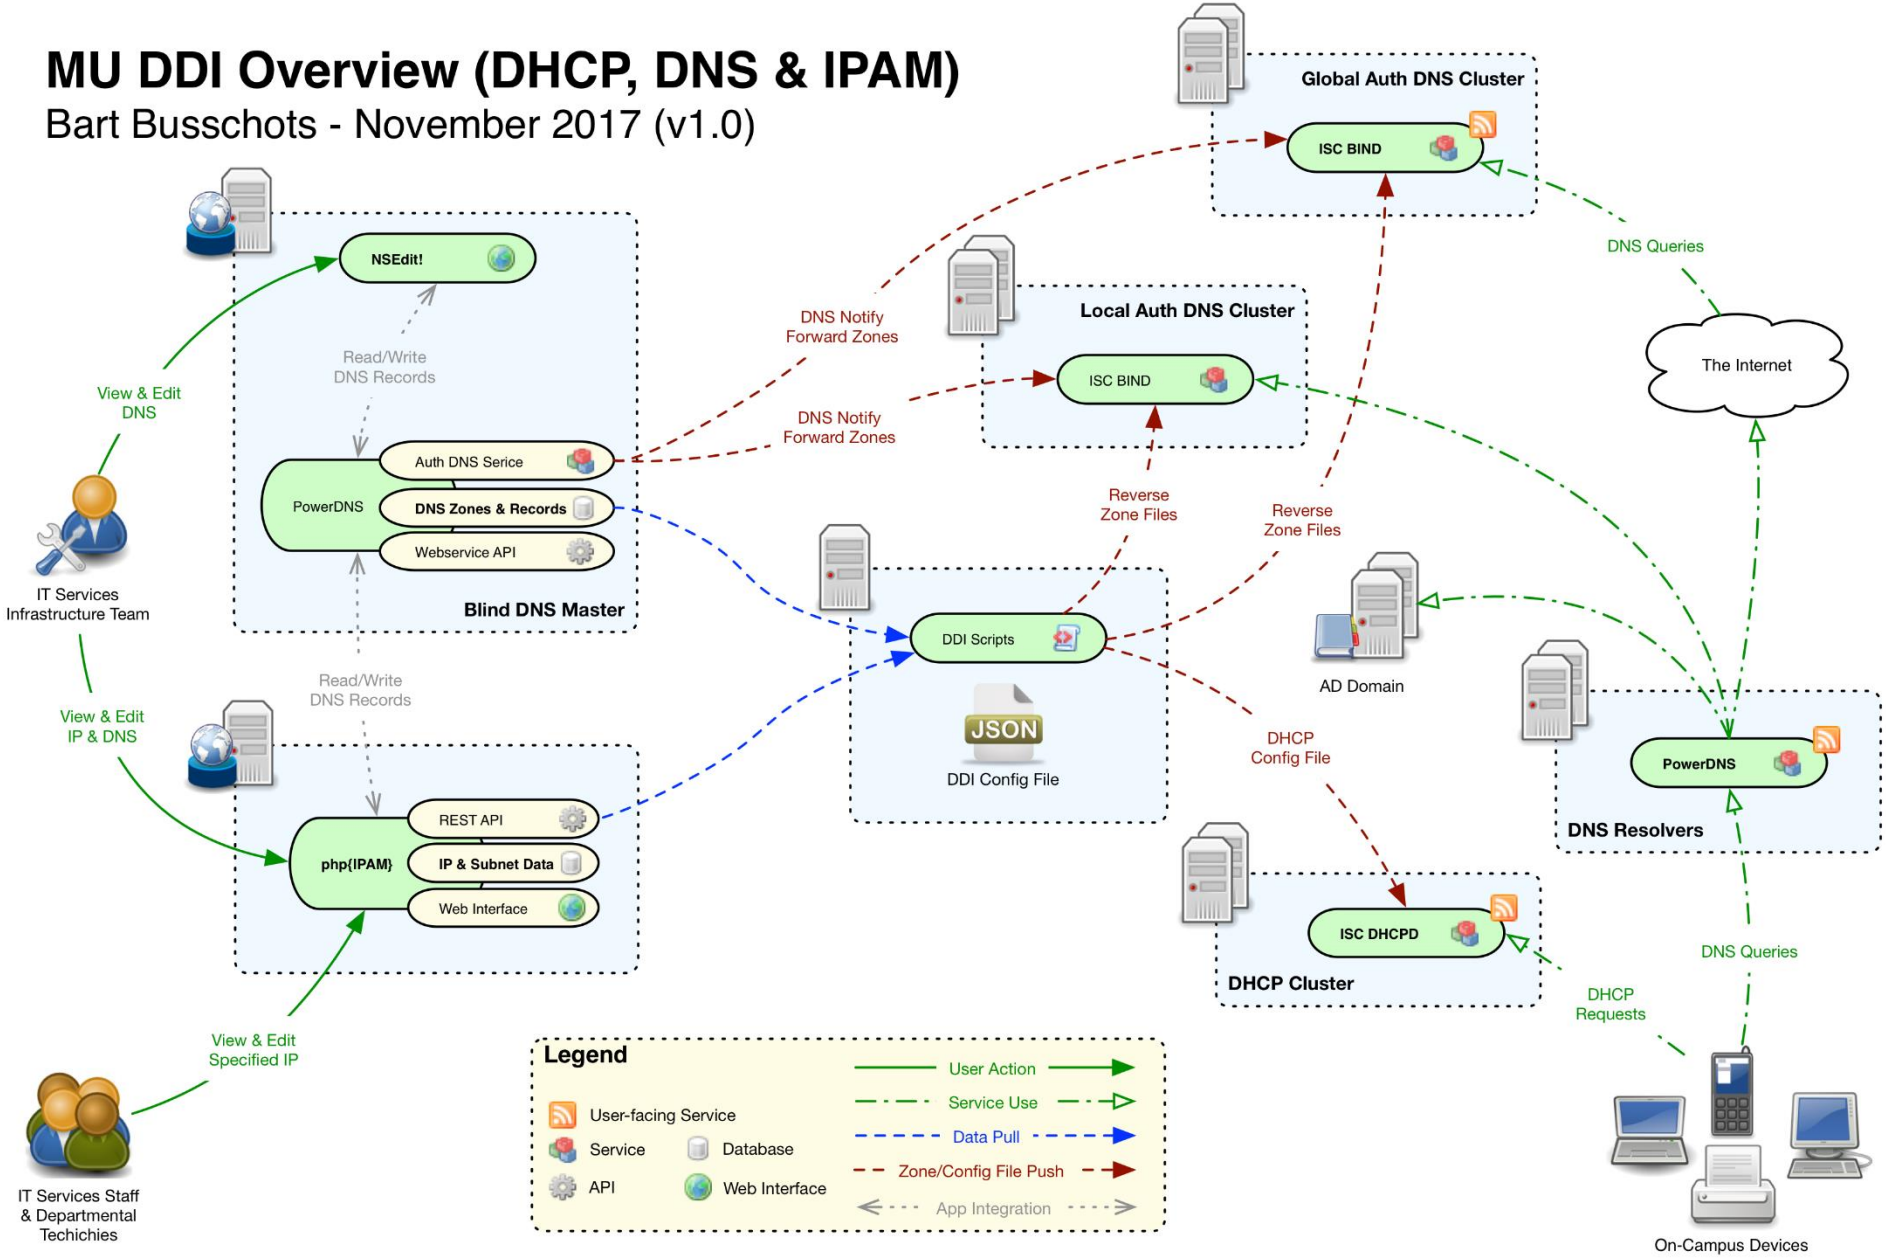

Maynooth University
National University of Ireland Maynooth

IT Services

Domain Name Changes

Before

- Single forward domain with two views:
  - Public (served by default)
  - Private (served to MU IPs)
- Single public reverse zone with two views:
  - Public (served by default)
  - Private (served to MU IPs)
- Private reverse zones served to MU IPs only
- No Dynamic DNS

After

- Multiple public forward zones
- Dedicated *private* forward zone
  - Stub served from 3rd party cloud provider (needed for TLS certificates)
  - True zone served from *Local* BIND cluster
- Two distinct versions of public reverse zone
  - Public version served from *Global* BIND cluster
  - Private version served from *Local* BIND cluster
- Private reverse zones served from *Local* BIND cluster
- AD providing dynamic DNS on delegated private sub-domain

php{IPAM} — Features

- **Flexible hierarchical organisation of IPv4 & IPv6 data**
  - Sections contain subnets and/or folders
  - Folders can contain other folders or subnets
  - Subnets contain other subnets and/or folders and information on IPs
- Multiple **authentication** options including **AD**, **LDAP** & **Radius**
- **Roles** can be assigned at the section, folder, and subnet level
- **Custom fields** can be defined for most data objects, including **subnets**
- Each IP entry must be assigned a **tag**. The tag list is customisable.
- Built-in **change tracking** on IP entries (who, when, what to what)
- Basic **PowerDNS integration**
  - Surfaces DNS records for tagged IPs in the php{IPAM} GUI
  - Allows DNS records be created for tagged IPs
  - Provides UI for managing DNS zones &records, but chokes on large zones
- **Poor data validation**

php{IPAM} — Our Configuration

- Authentication via AD
- Subnets organised into Sections by team, e.g. *Voice*, *BMS*, *Data Centre*, *Wifi Clients* & *PACRs*
- Custom Fields on Subnets:
  - **Notes:** a free-form text area
  - **Auto DHCP:** a yes/no toggle to indicate whether or not to include the subnet in the auto-generated DHCP config
  - **DHCP Options:** a free-form text area for specifying subnet-specific DHCP options, e.g. phone config options
  - **Deprecated:** a flag to indicate that the subnet is in the process of being decommissioned (no new devices should be added to it)
- Tags in use:
  - Standard tags: *Used* & *Reserved*
  - Custom tags: *DHCP Pool*, *DHCP Reservation* & *Managed by 3rd Party*

php{IPAM} UI — Section View

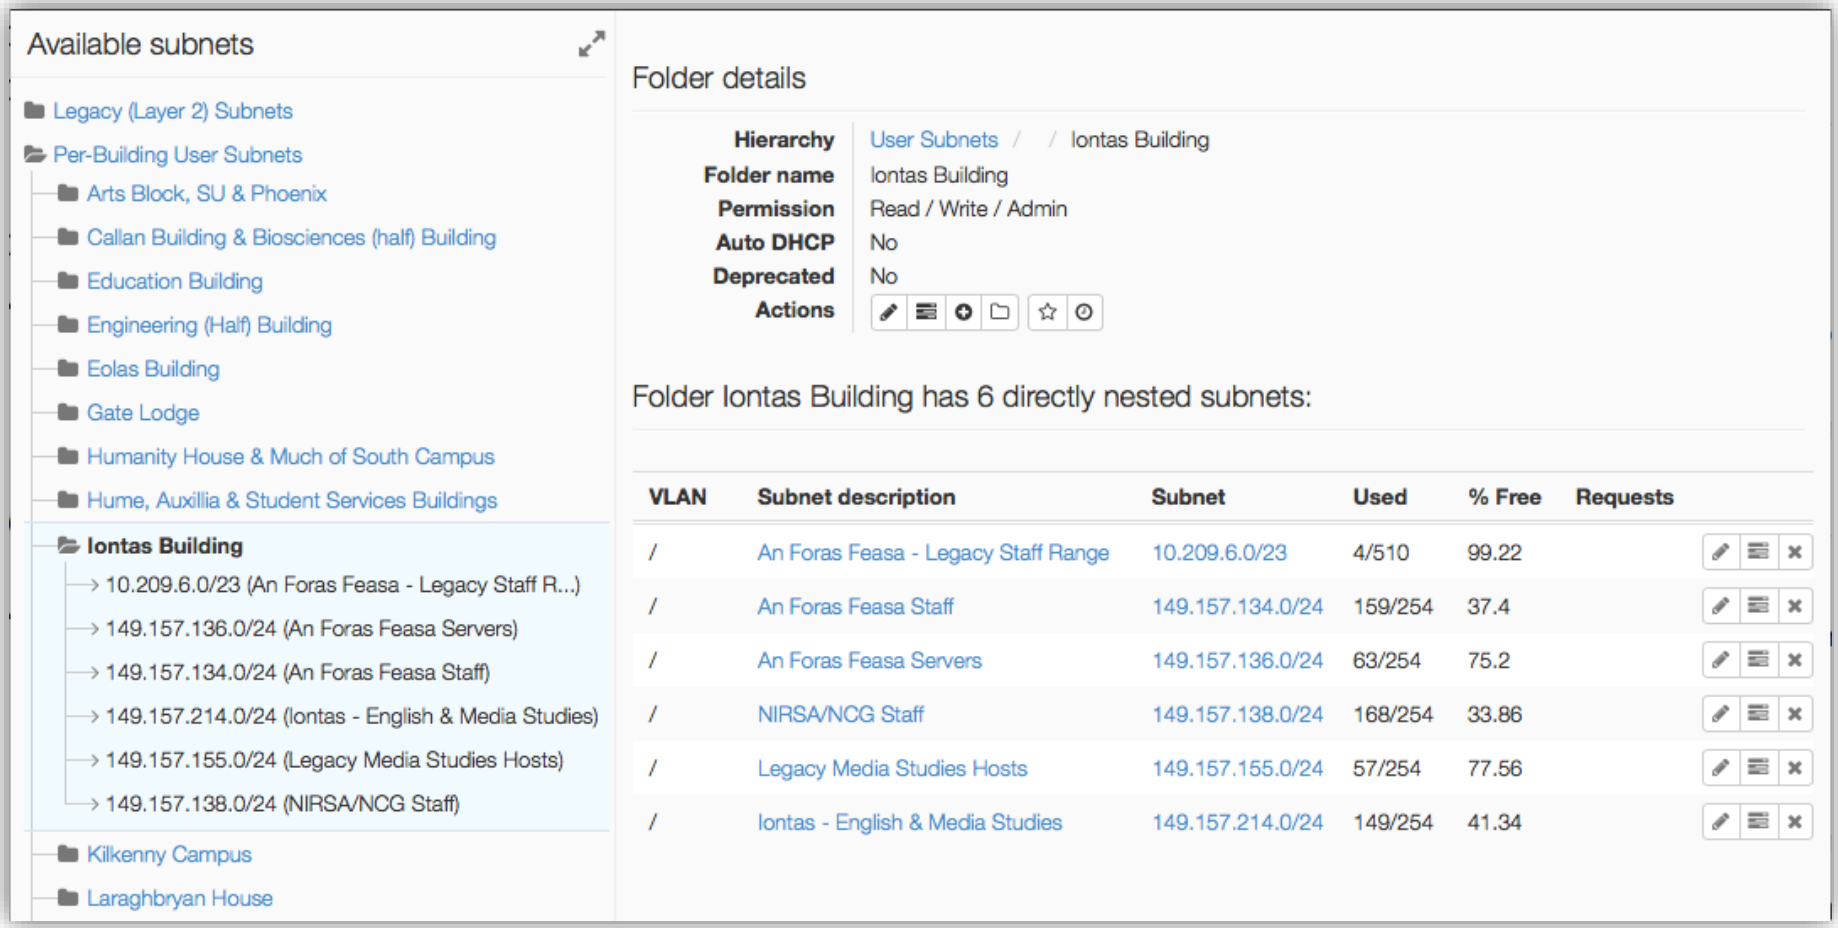

php{IPAM} UI — Subnet View

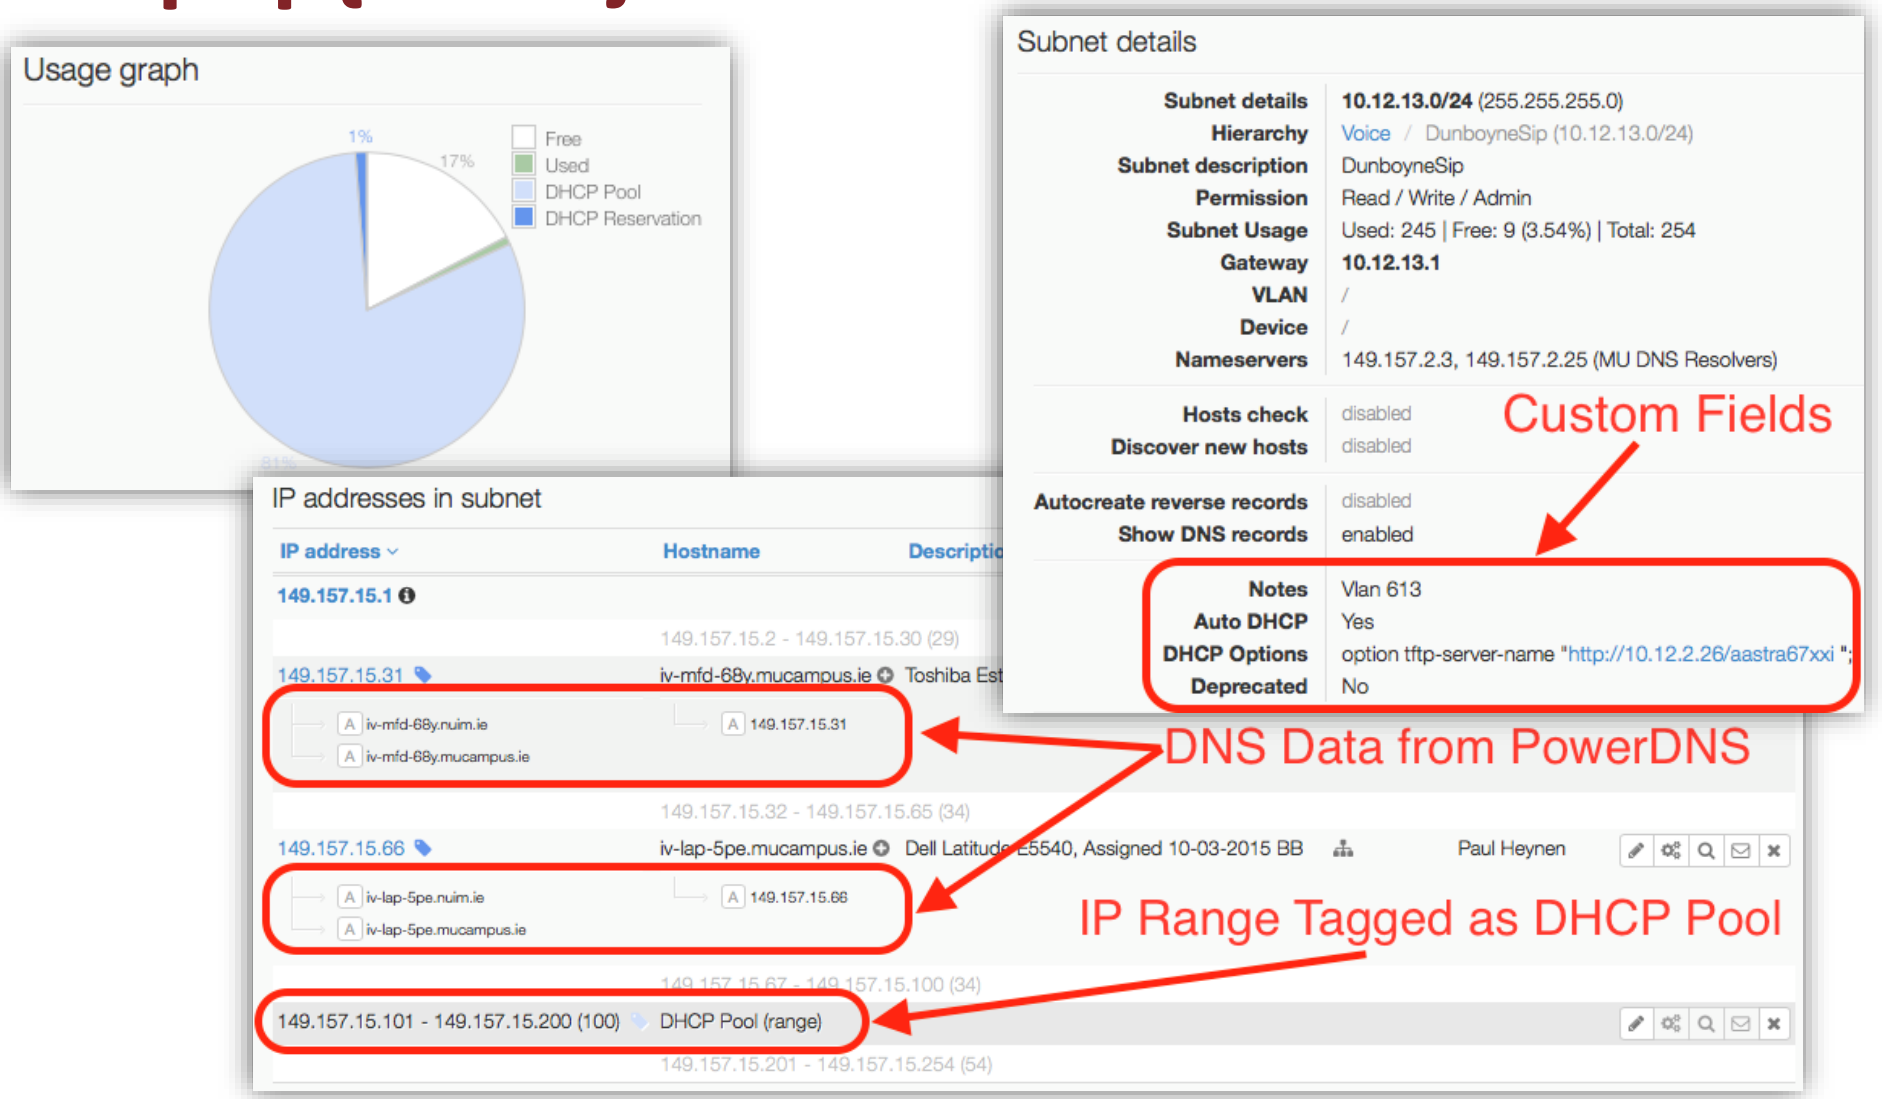

php{IPAM} UI — Subnet View

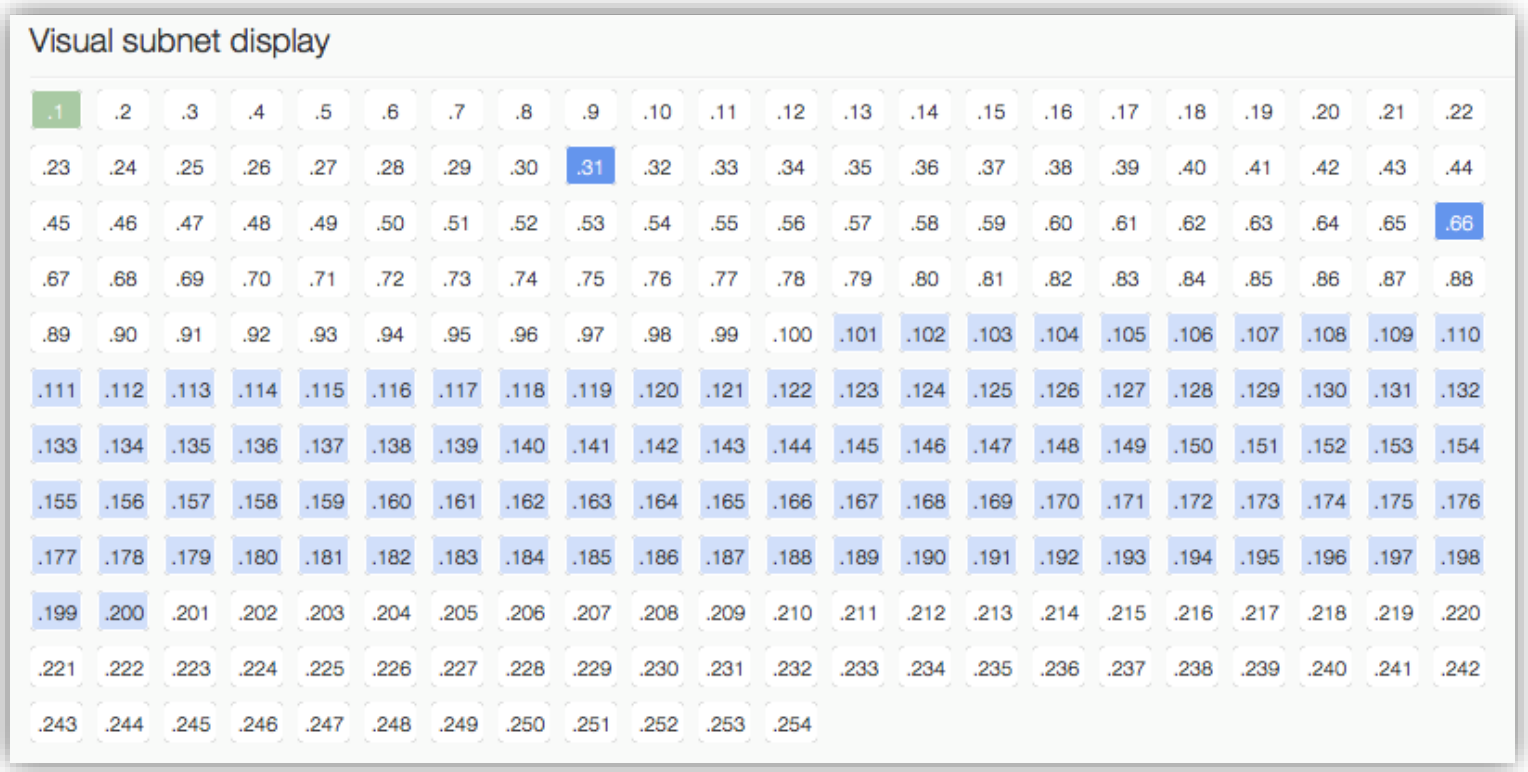

PowerDNS Suite (2 Distinct Products)

PowerDNS Recursor

- Features
  - Provides **DNS name resolution**
  - Light-weight & fast
  - Easy to configure
- Our Configuration
  - Cluster of two resolvers
  - Configured Forwarding Zones:
    - AD domain forwarded to DCs
    - Private MU domains forwarded to *Local* BIND cluster
    - MU reverse zones forwarded to *Local* BIND cluster

PowerDNS Server

- Features
  - Provides **hosting of DNS zones**
  - Supports multiple storage back-ends including RDBMSes
  - Compatible with BIND
  - Lightweight & fast
  - Easy to configure
- Our configuration
  - Single server
  - MySQL backend
  - *Blind master* to both *Global* and *Local* BIND clusters for all forward zones

PowerDNS Server GUIs

NSEdit!

- Free open-source PHP webapp
- View, edit & add zones and records
- UI works well for managing zones and adding records
- Good data validation when adding records
- UI works poorly for managing records with large domains — no search feature, just alphabetic list of records

MySQL Clients

- Table structure is very intuitive
- Generic MySQL client with search and filter features is often the quickest way to interact with PowerDNS

Other Open Source Components

ISC BIND

- *De-facto* Industry standard DNS server
- Robust & Full-featured
- Many years of experience with BIND in MU

ISC DHCPD

- *De-facto* Industry standard DHCP Server
- Robust & Full-featured
- Support for active-active clustering
- Many years of experience with DHCPD in MU

Our Scripts — “MU DDI Scripts”

- Middleware between the open-source components to implement MU business rules
  - Automatically generate & deploy DHCPD config
  - Automatically generate & deploy *global* and *local* versions of reverse DNS zones
  - Utility scripts
- Written in Perl
- Interaction with the open-source components through REST APIs and other officially supported and well documented mechanisms
- Actions and errors logged to central syslog server
- No *magic numbers* — all variables defined in single JSON-formatted config file (*the DDI Config file*)

DHCP Config Auto-Generation

- Generated config based on data in php{IPAM} & DDI config file
- Algorithm:
  1. Generate config file
  2. Validate generated config with `dhcpd -t -cf`
  3. Verify that both DHCP nodes are up with `dhcping`
  4. SFTP generated config to both nodes
  5. Re-start DHCPD on secondary, then primary via `systemd`
  6. Verify both nodes came back up with `dhcping`

DHCP Config Auto-Generation

- Generated config consists of three logical sections:
  1. Global settings (from DDI config file)
  2. Subnet declarations (from php{IPAM}, including our custom fields)
  3. Host declarations (from php{IPAM})
- Generated config heavily commented to make it human-readable for debugging purposes
  - Comments specify the source of the various pieces of information

DHCPD Config Snippet

```
#
# -- Subnet 149.157.7.192/26 (php{IPAM} id=224) --
#
# Description: Computer Science Servers
subnet 149.157.7.192 netmask 255.255.255.192{
  # -- Local Parameters --
  authoritative; # hard-coded
  option broadcast-address 149.157.7.255; # derived
  option routers 149.157.7.193; # from php{IPAM}

  # -- Custom Name Servers (from php{IPAM}) --
  option domain-name-servers 149.157.246.30;

  # -- Custom DHCP Options (from php{IPAM}) --
  option domain-search "cs.nuim.ie";
  option domain-name "cs.nuim.ie";

  # -- Dynamic IP Pool (from php{IPAM}) --
  pool{
    failover peer "mu_dhcp"; # from DDI config file
    range 149.157.7.247 149.157.7.254;
  }
}
```

Reverse DNS Zone Generation

- Multiple zone files need to be generated
- Generated zone files based on PowerDNS and DDI config file
- Algorithm (repeated for each needed zone file)
  1. Generate zone file
     - DDI config file defines which domains should be treated as *private*
     - DDI config file defines zone weightings for automatic conflict resolution between multiple `A` records for the same IP
  2. Validate generated file with `named-checkzone`
  3. SFTP generated file to *Local* or *Global* BIND master
  4. Trigger a zone reload via `rndc`

Utility Scripts

- Perform data integrity checks
  - Flag **single-field data validation issues** not dealt with by php{IPAM} (e.g. invalid MAC addresses or hostnames)
  - Flag **data integrity problems between records** (e.g. inconsistent mappings between MAC addresses and hostnames)
  - Flag **orphaned and obsolete records** (e.g. A records pointing to IPs within DHCP pools)
  - Flag inconsistencies between authoritative NS records for sub-domains and NS records returned by the delegated-to servers
- Bump the serial on all published domains
- Flush the caches on all resolvers

Maynooth University
National University of Ireland Maynooth

How did it Go?

The Big Change …

- Old and new infrastructure run in parallel (each DNS zone and subnet only served from one system at a time)
- Migration Process
  - DNS Resolvers deployed first
    - Legacy and new DHCP servers updated to point all clients at the new resolvers
    - DNS recursion disabled on legacy DNS (addressing auth+recurse security issue)
  - DHCP & Authoritative DNS Migrated in parallel
    - DHCP moved one subnet at a time by copying the data then changing the helper address on the router
    - Authoritative DNS moved one domain at a time (PowerDNS import from BIND zone file) then NS records and/or Forwarding zones updated

What's Improved?

- From our Staff's point of view:
  - Basic network connectivity 'just works' — no DHCP or DNS outages
  - Can move freely between buildings & departments with their devices
  - Can more easily use personally owned devices
  - Orders for University-owned devices processed more quickly
- From IT Services' point of view:
  - 'Single pane of glass' for all IPv4 & IPv6 network information
  - Reduction in work-load — very few devices need reservations, and more staff can now deal with the remaining few requests
  - Maintenance is easier because any single VM can be rebooted without triggering outages on any user-facing services
- From Departmental Technicians' Point of view:
  - Visibility into the configuration of and IP usage of the subnets relevant to them

Maynooth University
National University of Ireland Maynooth

Questions?